

#2303A51271

Batch:- 8

```
import numpy as np
import time
from numba import njit, prange

def serial_vector_op(A,B,alpha):
    C=np.empty_like(A)
    for i in range(len(A)):
        C[i]=alpha*A[i]+B[i]
    return C

@njit(parallel=True)
def parallel_vector_op(A,B,alpha):
    C=np.empty_like(A)
    for i in prange(len(A)):
        C[i]=alpha*A[i]+B[i]
    return C

def main():
    sizes=[10_000, 100_000,
    1_000_000, 2_000_000]
    alpha=2.5

    print(f"{'Size':>10} | {'Serial(s)':>10} | {'Parallel(s)':>12} | Speedup")
    print("-" * 55)

    for N in sizes:
        A = np.random.rand(N)
        B = np.random.rand(N)

        # Run once to compile Numba functions
        parallel_vector_op(A,B,alpha)

        #serial timing
        start=time.time()
        serial_vector_op(A,B,alpha)
        t_serial=time.time()-start

        #parallel timing
        start=time.time()
        parallel_vector_op(A,B,alpha)
        t_parallel=time.time()-start

        speedup=t_serial/t_parallel
        print(f"{{N:10d} | {t_serial:10.4f} | {t_parallel:12.4f} | {speedup:6.2f}}")
```

```

if __name__ == "__main__":
    main()

    Size | Serial(s) | Parallel(s) | Speedup
-----
    10000 |     0.0044 |     0.0001 | 81.47x
    100000 |    0.0463 |     0.0003 | 146.23x
    1000000 |   0.7449 |     0.0024 | 315.64x
2000000 |   1.7024 |     0.0047 | 366.00x

from itertools import starmap
import numpy as np
from numba import njit, prange
import time # Moved import time to the top

def serial_matmul(A,B):
    n = A.shape[0]
    C =
    np.zeros((n,n))
    for i
    in range(n):
        for j
        in range(n):
            for k
            in range(n):
                C[i,j] += A[i,k] * B[k,j]
    return C

@njit(parallel=True) # Corrected indentation
def parallel_outer_matmul(A,B):
    n =
    A.shape[0]
    C = np.zeros((n,n))
    for i in
    range (n):
        for j in range (n):
            for
    k in range(n):
                C[i,j] +=A[i,k]
    *B[k,j]
    return C # Corrected indentation

@njit(parallel=True) # Corrected indentation
def parallel_inner_matmul(A,B):
    n =
    A.shape[0]
    C = np.zeros((n,n))
    for idx
    in prange(n*n):
        i = idx // n
        j = idx
        % n
        for k in range(n):
            C[i,j] += A[i,k] * B [k,j]
    return C # Corrected indentation

def main():
    N = 300
    A = np.random.rand(N,N)
    B = np.random.rand(N,N)

```



```

# Warm-up calls for Numba JIT
parallel_outer_matmul(A,B)
parallel_inner_matmul(A,B)

    #serial timing
start = time.time()
serial_matmul(A,B)
    t_serial = time.time() - start

    #parallel outer      start
= time.time()
parallel_outer_matmul(A,B)
    t_parallel_outer = time.time() - start

    #parallel inner      start
= time.time()
parallel_inner_matmul(A,B)
    t_parallel_inner = time.time() - start

    print(f"Serial time      : {t_serial:.3f} s")
print(f"Parallel outer loop: {t_parallel_outer:.3f} s")
print(f"Collaped loop     : {t_parallel_inner:.3f} s")
print(f"Speedup(outer)    : {t_serial/t_parallel_outer:.2f} x")
print(f"Speedup(inner)    : {t_serial/t_parallel_inner:.2f} x") if
__name__ == "__main__":
    main()

Serial time      : 19.577 s
Parallel outer loop: 0.041 s
Collaped loop     : 0.024 s
Speedup(outer)    : 478.53 x
Speedup(inner)    : 813.28 x

from numba.np.ufunc import parallel
import numpy as np import time
import os
from numba import njit, prange, get_num_threads

def heavy_work(size):
    total=0.0    for i in
range(size):
    total+=(i%7)*0.5
    return total

def serial_processing(workloads):
    results=np.zeros(len(workloads))    for
    i in range(len(workloads)):
        results[i]=heavy_work(workloads[i])

```



```

    return results

@njit(parallel=True) def
parallel_processing(workloads):
results=np.zeros(len(workloads))
for i in prange(len(workloads)):
    total=0.0      for j in
range(workloads[i]):
total+=(j%7)*0.5
results[i]=total   return results

def main():      print("Load Balancing With Irregular
Workloads")      print("-----")
print(f"Logical CPU cores : {os.cpu_count()}")
print(f"Number of threads in use : {get_num_threads()}\n")
workloads=np.random.randint(10_000,500_000,size=40)
parallel_processing(workloads)

#serial timing
start=time.time()
    serial_processing(workloads)      t_serial=time.time()-start

#parallel timing
start=time.time()
    parallel_processing(workloads)      t_parallel=time.time()-start

    print(f"Serial execution time : {t_serial:.3f} s")
print(f"Parallel execution time : {t_parallel:.3f} s")
print(f"Speedup : {t_serial/t_parallel:.2f}x")

if __name__ == "__main__":
main()
Serial time       : 21.128 s
Parallel outer loop: 0.042 s
Collapsed loop    : 0.024 s
Speedup(outer)    : 508.36 x
Speedup(inner)    : 877.28 x
from numba.np.ufunc import parallel
import numpy as np import time
import os
from numba import njit,prange,get_num_threads

```



```

def heavy_work(size):
    total = 0.0    for i in
    range (size):      total +=
        (i % 7) * 0.5   return
    total

def serial_processing(workloads):
    results = np.zeros(len(workloads))    for
    i in range (len(workloads)):
        results[i] = heavy_work(workloads[i])
    return results

@njit (parallel=True) # Corrected 'true' to 'True'
def parallel_processing(workloads):    results =
    np.zeros(len(workloads))    for i in
    prange(len(workloads)):
        total= 0.0      for j in
        range (workloads[i]):
            total += (j % 7) * 0.5
        results[i] = total    return
    results

def main(): # Added colon here
    print("load balancing with irregular workloads")    print("-----"
-----)
    print(f"logical CPU cores : {os.cpu_count()}")
    print(f"NUmba threads : {get_num_threads()} \n")

    #sim
    workloads = np.random.randint(10_000,500_000,size=40) # Corrected
    'randit' to 'randint'

    #warm up
    parallel_processing(workloads)

    #serial timing      start =
    time.time()
    serial_processing(workloads)
    t_serial = time.time() - start

    #parallel      start
    = time.time()
    parallel_processing(workloads)
    t_parallel = time.time() - start

    print(f"Serial exe time: {t_serial:.3f} s")
    print(f"Parallel exe time: {t_parallel:.3f} s")
    print(f"Speedup: {t_serial/t_parallel:.2f} x")

```



```

if __name__ == "__main__": # Corrected indentation
main()

load balancing with irregular workloads
-----
logical CPU cores : 2
NUMba threads : 2

Serial exe time: 0.878 s
Parallel exe time: 0.019 s
Speedup: 46.47 x

import random
import time
import multiprocessing as mp

def chunk_sum(chunk):
    return sum(chunk)

def chunk_max(chunk):
    return max(chunk)

def parallel_reduce(data, func):
    cpu_count = mp.cpu_count()
    chunk_size = len(data) // cpu_count
    chunks = [data[i:i + chunk_size] for i in range(0, len(data), chunk_size)]

    with mp.Pool(cpu_count) as pool:
        results = pool.map(func, chunks)
    return func(results) if __name__ ==
"__main__":

    N = 10_000_000
    data = [random.randint(1, 100) for _ in range(N)]

    start = time.time()
    serial_sum = sum(data)
    serial_max = max(data)
    serial_time = time.time() - start

    start = time.time()
    parallel_sum = parallel_reduce(data, chunk_sum)
    parallel_max = parallel_reduce(data, chunk_max)
    parallel_time = time.time() - start

    print("Serial Sum:", serial_sum)
    print("Serial Max:", serial_max)
    print("Serial Time:", serial_time)

```



```

        print("Parallel Sum:", parallel_sum)
print("Parallel Max:", parallel_max)
print("Parallel Time:", parallel_time)

        print("Sum Correct:", serial_sum == parallel_sum)
print("Max Correct:", serial_max == parallel_max)
print("Speedup:", serial_time / parallel_time)

Serial Sum: 504913358
Serial Max: 100
Serial Time: 0.2057969570159912
Parallel Sum: 504913358
Parallel Max: 100
Parallel Time: 1.3743224143981934
Sum Correct: True
Max Correct: True
Speedup: 0.14974430661971583

import time import
numpy as np
from numba import njit, prange

# ----- Serial Version -----
@njit def
monte_carlo_pi_serial(n):
inside = 0      for i in
range(n):          x =
np.random.random()      y =
np.random.random()      if
x*x + y*y <= 1.0:
inside += 1      return 4.0 *
inside / n

# ----- Parallel Version -----
@njit(parallel=True) def monte_carlo_pi_parallel(n):      inside =
0      for i in prange(n):          x = np.random.random()      y
= np.random.random()      if x*x + y*y <= 1.0:
inside += 1      # Numba handles this as a safe reduction      return
4.0 * inside / n

# ----- Main Experiment -----def
main():      samples_list = [5_000_000, 10_000_000,
50_000_000]

```

```

# Warm-up for JIT compilation
monte_carlo_pi_serial(10_000)
monte_carlo_pi_parallel(10_000)

    print(f"{'Samples':>12} | {'Serial(s)':>10} | {'Parallel(s)':>12}
| {'Speedup':>8} | {'π (parallel)':>12}")
print("-" * 70)

    for n in samples_list:
        # Serial
t0 = time.time()
        pi_s = monte_carlo_pi_serial(n)
t1 = time.time()           serial_time =
t1 - t0

        # Parallel
t0 = time.time()
        pi_p = monte_carlo_pi_parallel(n)
t1 = time.time()           parallel_time =
t1 - t0

        speedup = serial_time / parallel_time

        print(f"{n:12d} | {serial_time:10.3f} | {parallel_time:12.3f}
| {speedup:8.2f} | {pi_p:12.6f}")

if __name__ == "__main__":
    main()

-----
```

Samples	Serial(s)	Parallel(s)	Speedup	π (parallel)
5000000	0.062	0.066	0.94	3.141834
10000000	0.221	0.251	0.88	3.142626
50000000	1.135	0.595	1.91	3.141647